

Assignment 3 - Part A (Utilizing AIR for Data Processing)

Name: Shivankar Pilligundla

Roll No: IMT2020016

Similar workflow as YSB is used to do the regex pattern match processing. Initially event generator generates a random 50 char pattern. Then its passed to an event filter within the same rank. This filters out the patterns based on regex and forwards the events to partial aggregator of same rank. Here window id's are calculated based on event time and count with in the rank is computed. Now these partial aggregators share their events with full aggregators of every possible rank(i.e all dataflows) and they unitedly perform final aggregation. Then these are forwarded to a collector at rank 0 that aggregates all the results and computes the avg latencies.

Results:

For the throughput of 100 the following latencies have been reported:

```
~/AIR/Release master !10 77 > mpirun -np 4 ./AIR YSB 100
*****AIR (c) 2020 Uni.lu*****
AIR INSTANCE AT RANK 2/4 | TP: 100 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 1000ms
AIR INSTANCE AT RANK 3/4 | TP: 100 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 1000ms
AIR INSTANCE AT RANK 4/4 | TP: 100 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 1000ms
AIR INSTANCE AT RANK 1/4 | TP: 100 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 1000ms
WID: 0 count: 144 latency: 1001 SIZE 0 CAP 2400

#1 COUNT: 144 AVG_LATENCY: 1001 N=1
WID: 1 count: 384 latency: 1 SIZE 0 CAP 2400

#2 COUNT: 384 AVG_LATENCY: 501 N=1
WID: 2 count: 240 latency: 1 SIZE 0 CAP 2400

#3 COUNT: 240 AVG_LATENCY: 334 N=1
WID: 3 count: 384 latency: 1001 SIZE 0 CAP 2400

#4 COUNT: 384 AVG_LATENCY: 501 N=1
WID: 4 count: 288 latency: 1001 SIZE 0 CAP 2400

#5 COUNT: 288 AVG_LATENCY: 601 N=1
WID: 5 count: 96 latency: 1 SIZE 0 CAP 2400

#6 COUNT: 96 AVG_LATENCY: 501 N=1
WID: 6 count: 288 latency: 2001 SIZE 0 CAP 2400

#7 COUNT: 288 AVG_LATENCY: 715 N=1
WID: 7 count: 48 latency: 6001 SIZE 0 CAP 2400

#8 COUNT: 48 AVG_LATENCY: 1376 N=1
WID: 8 count: 336 latency: 2001 SIZE 0 CAP 2400

#9 COUNT: 336 AVG_LATENCY: 1445 N=1
WID: 9 count: 240 latency: 1001 SIZE 0 CAP 2400

#10 COUNT: 240 AVG_LATENCY: 1401 N=1
WID: 10 count: 240 latency: 1 SIZE 0 CAP 2400

#11 COUNT: 240 AVG_LATENCY: 1273 N=1
WID: 11 count: 336 latency: 1001 SIZE 0 CAP 2400

#12 COUNT: 336 AVG_LATENCY: 1251 N=1
WID: 12 count: 384 latency: 1001 SIZE 0 CAP 2400

#13 COUNT: 384 AVG_LATENCY: 1231 N=1
WID: 13 count: 288 latency: 1 SIZE 0 CAP 2400

#14 COUNT: 288 AVG_LATENCY: 1143 N=1
WID: 14 count: 288 latency: 1 SIZE 0 CAP 2400

#15 COUNT: 288 AVG_LATENCY: 1067 N=1
WID: 15 count: 480 latency: 1 SIZE 0 CAP 2400

#16 COUNT: 480 AVG_LATENCY: 1001 N=1
```

When the throughput is increased to 10000:

```
~/AIR/Release master !10 ?? > mpirun -np 4 ./AIR YSB 10000
*****AIR (c) 2020 Uni.lu*****
AIR INSTANCE AT RANK 1/4 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 2/4 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 3/4 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 4/4 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
WID: 0 count: 35000 latency: 4 SIZE 0 CAP 2400

#1 COUNT: 35000 AVG_LATENCY: 4 N=1
WID: 1 count: 40000 latency: 2002 SIZE 0 CAP 2400
#2 COUNT: 40000 AVG_LATENCY: 1003 N=1
WID: 2 count: 40000 latency: 4 SIZE 0 CAP 2400
#3 COUNT: 40000 AVG_LATENCY: 670 N=1
WID: 3 count: 45000 latency: 1002 SIZE 0 CAP 2400
#4 COUNT: 45000 AVG_LATENCY: 753 N=1
WID: 4 count: 35000 latency: 2002 SIZE 0 CAP 2400
#5 COUNT: 35000 AVG_LATENCY: 1002 N=1
WID: 5 count: 40000 latency: 4 SIZE 0 CAP 2400
#6 COUNT: 40000 AVG_LATENCY: 836 N=1
WID: 6 count: 30000 latency: 4 SIZE 0 CAP 2400
#7 COUNT: 30000 AVG_LATENCY: 717 N=1
WID: 7 count: 40000 latency: 1002 SIZE 0 CAP 2400
#8 COUNT: 40000 AVG_LATENCY: 753 N=1
WID: 8 count: 45000 latency: 4 SIZE 0 CAP 2400
#9 COUNT: 45000 AVG_LATENCY: 669 N=1
WID: 9 count: 30000 latency: 2002 SIZE 0 CAP 2400
#10 COUNT: 30000 AVG_LATENCY: 803 N=1
WID: 10 count: 35000 latency: 1002 SIZE 0 CAP 2400
#11 COUNT: 35000 AVG_LATENCY: 821 N=1
WID: 11 count: 35000 latency: 4 SIZE 0 CAP 2400
#12 COUNT: 35000 AVG_LATENCY: 753 N=1
WID: 12 count: 30000 latency: 4 SIZE 0 CAP 2400
#13 COUNT: 30000 AVG_LATENCY: 695 N=1
WID: 13 count: 25000 latency: 4 SIZE 0 CAP 2400
#14 COUNT: 25000 AVG_LATENCY: 646 N=1
WID: 14 count: 25000 latency: 4 SIZE 0 CAP 2400
#15 COUNT: 25000 AVG_LATENCY: 603 N=1
WID: 15 count: 30000 latency: 4 SIZE 0 CAP 2400
#16 COUNT: 30000 AVG_LATENCY: 565 N=1
```

Throughput(50k) - this was the max:

```
~/AIR/Release master !10 ?? > mpirun -np 4 ./AIR YSB 50000
*****AIR (c) 2020 Uni.lu*****
AIR INSTANCE AT RANK 4/4 | TP: 50000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 2/4 | TP: 50000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 1/4 | TP: 50000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 3/4 | TP: 50000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
WID: 0 count: 225000 latency: 32 SIZE 0 CAP 2400

#1 COUNT: 225000 AVG_LATENCY: 32 N=1
WID: 1 count: 75000 latency: 15 SIZE 0 CAP 2400
#2 COUNT: 75000 AVG_LATENCY: 23 N=1
WID: 2 count: 175000 latency: 1006 SIZE 0 CAP 2400
#3 COUNT: 175000 AVG_LATENCY: 351 N=1
WID: 3 count: 200000 latency: 16 SIZE 0 CAP 2400
#4 COUNT: 200000 AVG_LATENCY: 267 N=1
WID: 4 count: 175000 latency: 16 SIZE 0 CAP 2400
#5 COUNT: 175000 AVG_LATENCY: 217 N=1
WID: 5 count: 125000 latency: 3007 SIZE 0 CAP 2400
#6 COUNT: 125000 AVG_LATENCY: 682 N=1
WID: 6 count: 125000 latency: 1007 SIZE 0 CAP 2400
#7 COUNT: 125000 AVG_LATENCY: 728 N=1
WID: 7 count: 175000 latency: 25 SIZE 0 CAP 2400
#8 COUNT: 175000 AVG_LATENCY: 640 N=1
WID: 8 count: 125000 latency: 1007 SIZE 0 CAP 2400
#9 COUNT: 125000 AVG_LATENCY: 681 N=1
WID: 9 count: 200000 latency: 16 SIZE 0 CAP 2400
#10 COUNT: 200000 AVG_LATENCY: 614 N=1
WID: 10 count: 150000 latency: 2011 SIZE 0 CAP 2400
#11 COUNT: 150000 AVG_LATENCY: 741 N=1
WID: 11 count: 100000 latency: 1007 SIZE 0 CAP 2400
#12 COUNT: 100000 AVG_LATENCY: 763 N=1
WID: 12 count: 150000 latency: 1007 SIZE 0 CAP 2400
#13 COUNT: 150000 AVG_LATENCY: 782 N=1
WID: 13 count: 100000 latency: 16 SIZE 0 CAP 2400
#14 COUNT: 100000 AVG_LATENCY: 727 N=1
WID: 14 count: 125000 latency: 16 SIZE 0 CAP 2400
#15 COUNT: 125000 AVG_LATENCY: 680 N=1
WID: 15 count: 200000 latency: 27 SIZE 0 CAP 2400
#16 COUNT: 200000 AVG_LATENCY: 639 N=1
```

Experimenting with different dataflows. The max throughput supported by a single dataflow was 10k. So I've decided to benchmark different no of dataflows using 10k as throughput:

1 dataflow:

```
[~/AIR/Release master !10 ?7 > mpirun -np 1 ./AIR YSB 10000  
*****AIR (c) 2020 Uni.lu*****  
AIR INSTANCE AT RANK 1/1 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms  
WID: 0 count: 25000 latency: 12 SIZE 0 CAP 2400  
#1 COUNT: 25000 AVG_LATENCY: 12 N=1  
WID: 1 count: 10000 latency: 8005 SIZE 0 CAP 2400  
#2 COUNT: 10000 AVG_LATENCY: 4008 N=1  
WID: 2 count: 35000 latency: 11 SIZE 0 CAP 2400  
#3 COUNT: 35000 AVG_LATENCY: 2676 N=1  
WID: 3 count: 25000 latency: 19 SIZE 0 CAP 2400  
#4 COUNT: 25000 AVG_LATENCY: 2011 N=1  
WID: 4 count: 50000 latency: 11 SIZE 0 CAP 2400  
#5 COUNT: 50000 AVG_LATENCY: 1611 N=1  
WID: 5 count: 35000 latency: 12 SIZE 0 CAP 2400  
#6 COUNT: 35000 AVG_LATENCY: 1345 N=1  
WID: 6 count: 30000 latency: 12 SIZE 0 CAP 2400  
#7 COUNT: 30000 AVG_LATENCY: 1154 N=1  
WID: 7 count: 50000 latency: 12 SIZE 0 CAP 2400  
#8 COUNT: 50000 AVG_LATENCY: 1011 N=1  
WID: 8 count: 25000 latency: 12 SIZE 0 CAP 2400  
#9 COUNT: 25000 AVG_LATENCY: 900 N=1  
WID: 9 count: 35000 latency: 1005 SIZE 0 CAP 2400  
#10 COUNT: 35000 AVG_LATENCY: 911 N=1
```

5 dataflows:

```
[~/AIR/Release master !10 ?7 > mpirun -np 5 ./AIR YSB 10000
```

```
*****AIR (c) 2020 Uni.lu*****
```

```
AIR INSTANCE AT RANK 1/5 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 2/5 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 3/5 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 4/5 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 5/5 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
WID: 0 count: 25000      latency: 5 SIZE 0 CAP 2400
```

```
#1 COUNT: 25000      AVG_LATENCY: 5      N=1
```

```
WID: 1 count: 30000      latency: 3003 SIZE 0 CAP 2400
```

```
#2 COUNT: 30000      AVG_LATENCY: 1504      N=1
```

```
WID: 2 count: 30000      latency: 3003 SIZE 0 CAP 2400
```

```
#3 COUNT: 30000      AVG_LATENCY: 2003      N=1
```

```
WID: 3 count: 15000      latency: 4003 SIZE 0 CAP 2400
```

```
#4 COUNT: 15000      AVG_LATENCY: 2503      N=1
```

```
WID: 4 count: 15000      latency: 4003 SIZE 0 CAP 2400
```

```
#5 COUNT: 15000      AVG_LATENCY: 2803      N=1
```

```
WID: 5 count: 25000      latency: 3003 SIZE 0 CAP 2400
```

```
#6 COUNT: 25000      AVG_LATENCY: 2836      N=1
```

```
WID: 6 count: 55000      latency: 5 SIZE 0 CAP 2400
```

```
#7 COUNT: 55000      AVG_LATENCY: 2432      N=1
```

```
WID: 7 count: 25000      latency: 2003 SIZE 0 CAP 2400
```

```
#8 COUNT: 25000      AVG_LATENCY: 2378      N=1
```

```
WID: 8 count: 30000      latency: 5 SIZE 0 CAP 2400
```

```
#9 COUNT: 30000      AVG_LATENCY: 2114      N=1
```

```
WID: 9 count: 20000      latency: 3003 SIZE 0 CAP 2400
```

```
#10 COUNT: 20000      AVG_LATENCY: 2203      N=1
```


10 dataflows:

```
[~/AIR/Release master !10 ?7 > mpirun -np 10 ./AIR YSB 10000

*****AIR (c) 2020 Uni.lu*****

AIR INSTANCE AT RANK 8/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 9/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 1/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 2/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 3/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 4/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 5/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 6/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 7/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
AIR INSTANCE AT RANK 10/10 | TP: 10000 | MSG/SEC/RANK: 2 | AGGR_WINDOW: 10000ms
WID: 0 count: 20000 latency: 1014 SIZE 0 CAP 2400

#1 COUNT: 20000 AVG_LATENCY: 1014 N=1
WID: 1 count: 25000 latency: 2032 SIZE 0 CAP 2400

#2 COUNT: 25000 AVG_LATENCY: 1523 N=1
WID: 2 count: 30000 latency: 9 SIZE 0 CAP 2400

#3 COUNT: 30000 AVG_LATENCY: 1018 N=1
WID: 3 count: 30000 latency: 37 SIZE 0 CAP 2400

#4 COUNT: 30000 AVG_LATENCY: 773 N=1
WID: 4 count: 25000 latency: 12 SIZE 0 CAP 2400

#5 COUNT: 25000 AVG_LATENCY: 620 N=1
WID: 5 count: 30000 latency: 8 SIZE 0 CAP 2400

#6 COUNT: 30000 AVG_LATENCY: 518 N=1
WID: 6 count: 40000 latency: 6 SIZE 0 CAP 2400

#7 COUNT: 40000 AVG_LATENCY: 445 N=1
WID: 7 count: 35000 latency: 8 SIZE 0 CAP 2400

#8 COUNT: 35000 AVG_LATENCY: 390 N=1
WID: 8 count: 40000 latency: 1066 SIZE 0 CAP 2400

#9 COUNT: 40000 AVG_LATENCY: 465 N=1
WID: 9 count: 35000 latency: 8 SIZE 0 CAP 2400

#10 COUNT: 35000 AVG_LATENCY: 420 N=1
WID: 10 count: 25000 latency: 1014 SIZE 0 CAP 2400
```